

```
#import libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

#Load the data
from google.colab import files
uploaded = files.upload()
df = pd.read_csv('Breast_Cancer_Database.csv')
df.head(7)
```

📁 Choose Files Breast_Can...abase.csv

- **Breast_Cancer_Database.csv**(application/vnd.ms-excel) - 127838 bytes, last modified: 7/14/2020 - 100% done
- Saving Breast_Cancer_Database.csv to Breast_Cancer_Database (2).csv

	ID	Diagnosis	Radius mean	Texture_Mean	Perimeter_Mean	Area_Mean	Smoothness_Me
0	842302	M	17.99	10.38	122.80	1001.0	0.118
1	842517	M	20.57	17.77	132.90	1326.0	0.084
2	843009	M	19.69	21.25	130.00	1203.0	0.109
3	843483	M	11.42	20.38	77.58	386.1	0.142
4	843584	M	20.29	14.34	135.10	1297.0	0.100
5	843786	M	12.45	15.70	82.57	477.1	0.127
6	844359	M	18.25	19.98	119.60	1040.0	0.094

```
#Count the number of rows and columns in the data set
df.shape
```

📁 (569, 33)

```
#Count the number of empty(NaN,NAN,na,0) values in each column
df.isna().sum()
```

📁

```

ID                0
Diagnosis          0
Radius_mean       0
Texture_Mean      0
Perimeter_Mean    0
Area_Mean         0
Smoothness_Mean   0
Compactness_Mean  0
Concavity_Mean    0
Concave_Points_Mean 0
Symmetry_Mean     0
Fractal_Dimension_Mean 0
Radius_se        0
Texture_se       0
Perimeter_se     0
Area_se          0
Smoothness_se    0
Compactness_se   0
Concavity_se     0
Concave points_se 0
Symmetry_se      0

```

#Drop the column with all missing values

```
df = df.dropna(axis=1)
```

```

Perimeter_worst      0

```

#Get the new count of the number of rows and columns

```
df.shape
```

```
(569, 32)
```

```

Symmetry_Worst      0

```

#Get a count of the number of Malignant(M) or Benign(B) cells

```
df['Diagnosis'].value_counts()
```

```


B      357
M      212
Name: Diagnosis, dtype: int64

```

#Visualize the count

```
sns.countplot(df['Diagnosis'],label='count')
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f1a158a2748>
```



```
#Look at the data types to see which columns need to be encoded.
df.dtypes
```

```
↳ ID          int64
   Diagnosis   object
   Radius_mean float64
   Texture_Mean float64
   Perimeter_Mean float64
   Area_Mean   float64
   Smoothness_Mean float64
   Compactness_Mean float64
   Concavity_Mean float64
   Concave_Points_Mean float64
   Symmetry_Mean float64
   Fractal_Dimension_Mean float64
   Radius_se   float64
   Texture_se  float64
   Perimeter_se float64
   Area_se     float64
   Smoothness_se float64
   Compactness_se float64
   Concavity_se float64
   Concave points_se float64
   Symmetry_se float64
   Fractal_Dimension_se float64
   Radius_worst float64
   Texture_worst float64
   Perimeter_worst float64
   Area_Worst   float64
   Smoothness_Worst float64
   Compactness_Worst float64
   Concavity_Worst float64
   Concave_Points_Worst float64
   Symmetry_Worst float64
   Fractal_Dimension_Worst float64
   dtype: object
```

```
#Encode the categorical data values
```

```
from sklearn.preprocessing import LabelEncoder
```

```
labelencoder_Y = LabelEncoder()
```


```
df.iloc[:,1] = labelencoder_Y.fit_transform(df.iloc[:,1].values)
```

```
#Create a pair plot
```

```
sns.pairplot(df.iloc[:,1:5], hue='Diagnosis')
```

```
↳
```

<seaborn.axisgrid.PairGrid at 0x7f1a0dd29e48>



```
#Print the 1st 5 rows of the new data
df.head(5)
```



	ID	Diagnosis	Radius_mean	Texture_Mean	Perimeter_Mean	Area_Mean	Smoothness_Mean
0	187	1	17.99	10.38	122.80	1001.0	0.11840
1	189	1	20.57	17.77	132.90	1326.0	0.08474
2	193	1	19.69	21.25	130.00	1203.0	0.10960
3	194	1	11.42	20.38	77.58	386.1	0.14250
4	195	1	20.29	14.34	135.10	1297.0	0.10030

```
#Get the correlation of the columns
df.iloc[:,1:12].corr()
```



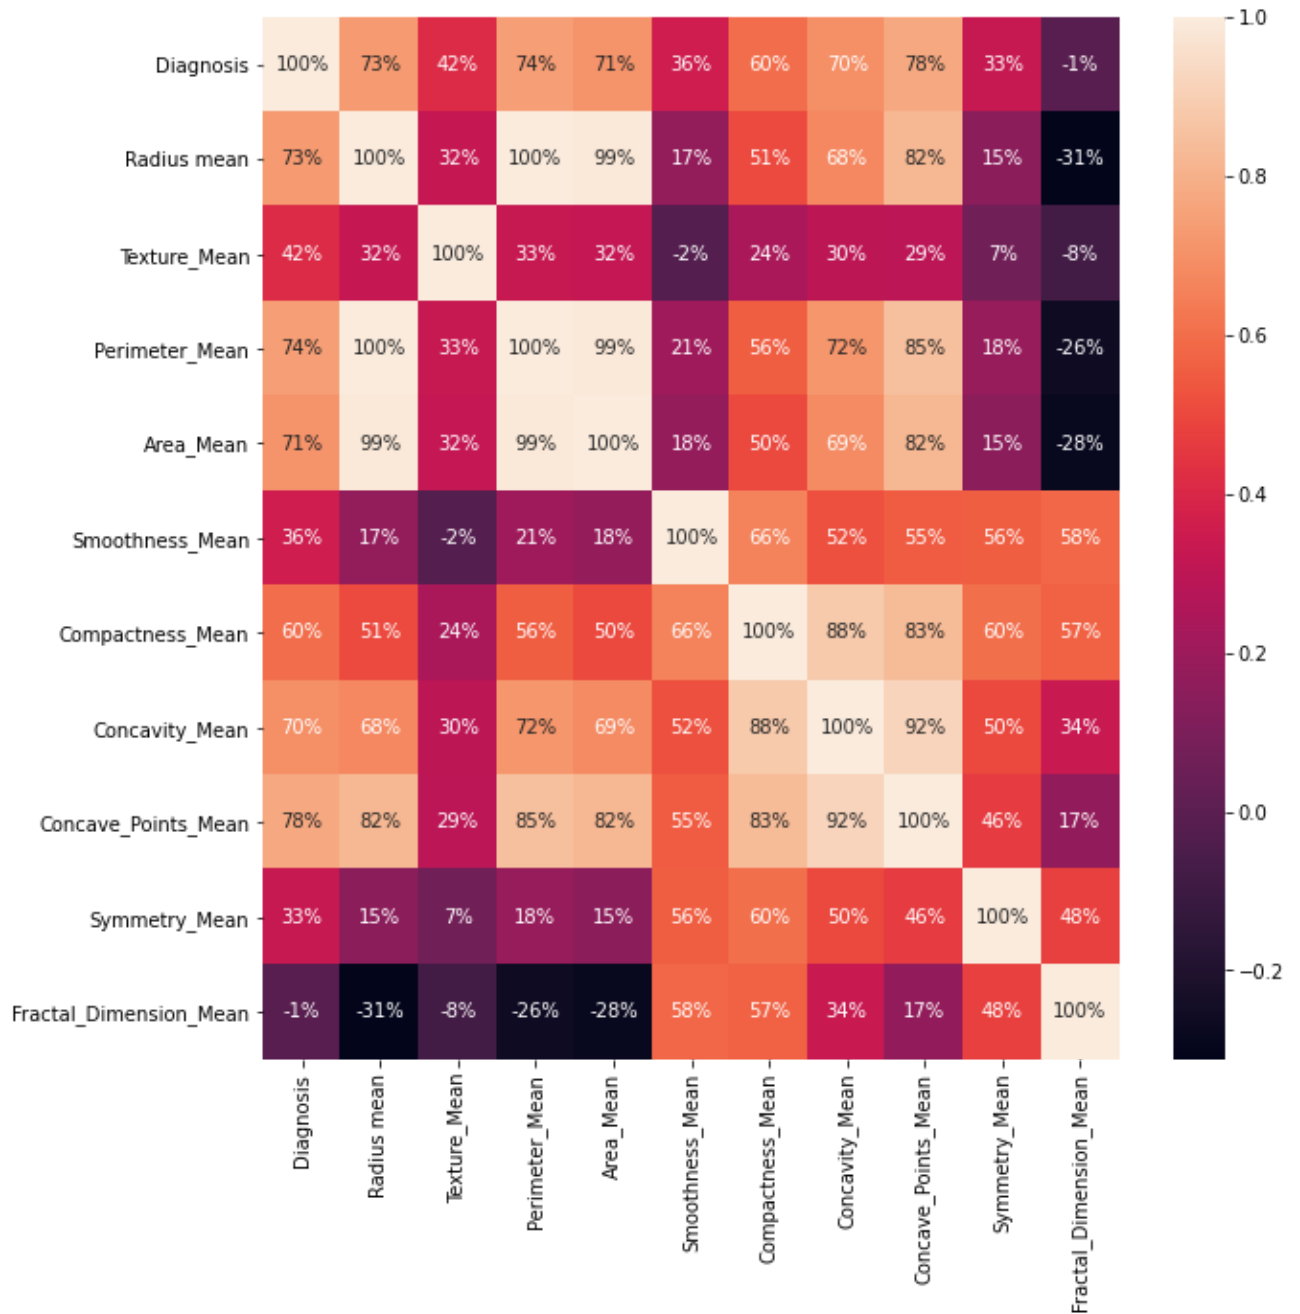
	Diagnosis	Radius mean	Texture_Mean	Perimeter_Mean	Area_Mean
Diagnosis	1.000000	0.730029	0.415185	0.742636	0.708984
Radius mean	0.730029	1.000000	0.323782	0.997855	0.987357

#Visualize the correlation

```
plt.figure(figsize=(10,10))
```

```
sns.heatmap(df.iloc[:,1:12].corr(), annot=True, fmt='.0%')
```

↳ <matplotlib.axes._subplots.AxesSubplot at 0x7f1a0b9ec0f0>



#Split the data set into independent (x) and dependent (y) data sets

```
x = df.iloc[:,2:31].values
```

```
y = df.iloc[:,1].values
```

#Split the data set into 75% training and 25% testing

```

#Split the data set into 75% training and 25% testing
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.25 , random_state

#Scale the data (Feature scaling)
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
x_train = sc.fit_transform(x_train)
x_test = sc.fit_transform(x_test)

#Create a function for the models
def models(x_train, y_train):

    #Logistic Regression
    from sklearn.linear_model import LogisticRegression
    log = LogisticRegression(random_state=0)
    log.fit(x_train, y_train)

    #Decision Tree
    from sklearn.tree import DecisionTreeClassifier
    tree = DecisionTreeClassifier(criterion = 'entropy', random_state=0)
    tree.fit(x_train, y_train)

    #Random Forest Classifier
    from sklearn.ensemble import RandomForestClassifier
    forest = RandomForestClassifier(n_estimators = 10, criterion = 'entropy', random_state =
    forest.fit(x_train, y_train)

    #Print the models accuracy on the training data
    print('[0]Logistic Regression Training Accuracy:', log.score(x_train, y_train))
    print('[1]Decision Tree Classifier Training Accuracy:', tree.score(x_train, y_train))
    print('[2]Random Forest Classifier Training Accuracy:', forest.score(x_train, y_train))

    return log, tree, forest

#Getting all of the models
model = models(x_train, y_train)

[0]Logistic Regression Training Accuracy: 0.9906103286384976
[1]Decision Tree Classifier Training Accuracy: 1.0
[2]Random Forest Classifier Training Accuracy: 0.9953051643192489

#Test model accuracy on data on confusion matrix
from sklearn.metrics import confusion_matrix

for i in range( len(model) ):
    print('Model ', i)
    cm = confusion_matrix(y_test, model[i].predict(x_test))

    TP = cm[0][0]
    TN = cm[1][1]

```

```
FN = cm[1][0]
FP = cm[0][1]

print(cm)
print('Testing Accuracy = ', (TP + TN)/(TP + TN + FN +FP))
print()
```

```
↳ Model 0
[[86  4]
 [ 3 50]]
Testing Accuracy = 0.951048951048951

Model 1
[[83  7]
 [ 2 51]]
Testing Accuracy = 0.9370629370629371

Model 2
[[87  3]
 [ 2 51]]
Testing Accuracy = 0.965034965034965
```

```
#Show another way to get metrics of the models
from sklearn.metrics import classification_report
from sklearn.metrics import accuracy_score

for i in range( len(model) ):
    print('Model ', i)
    print( classification_report(y_test, model[i].predict(x_test)))
    print( accuracy_score(y_test, model[i].predict(x_test)))
    print()
```

↳

Model	0				
		precision	recall	f1-score	support
	0	0.97	0.96	0.96	90
	1	0.93	0.94	0.93	53
	accuracy			0.95	143
	macro avg	0.95	0.95	0.95	143
	weighted avg	0.95	0.95	0.95	143

0.951048951048951

```
#Print the prediction of Random Forest Classifier Model
pred = model[2].predict(x_test)
print(pred)
print(y_test)
```

[1 0 0 0 0 0 0 0 0 0 0 1 0 0 1 1 1 0 1 1 1 1 1 0 0 1 0 0 1 0 1 0 1 0 1 0 1 0 1 0
1 0 1 0 0 1 0 0 1 0 0 0 1 1 1 1 0 0 0 0 0 0 0 1 1 1 0 0 1 0 1 1 1 0 0 1 0 0
1 0 0 0 0 0 1 1 1 0 1 0 0 0 1 1 0 1 0 1 0 0 1 0 0 0 0 0 0 0 0 1 0 1 0 1 0 1 1 0
1 1 0 0 0 0 0 0 0 0 0 0 1 0 1 0 0 0 0 0 1 0 0 0 0 0 0 0 1 1 0 0 0 1]
[1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 1 1 1 1 1 1 0 0 1 0 0 1 0 1 0 1 0 1 0 1 0
1 0 1 1 0 1 0 0 1 0 0 0 1 1 1 1 0 0 0 0 0 0 1 1 1 0 0 1 0 1 1 1 0 0 1 0 1
1 0 0 0 0 0 1 1 1 0 1 0 0 0 1 1 0 1 0 1 0 0 1 0 0 0 0 0 0 0 1 0 1 0 1 0 1 1 0
1 1 0 0 0 0 0 0 0 0 0 0 1 0 1 0 0 0 0 0 1 0 0 0 0 0 0 0 1 1 0 0 0 1]

		precision	recall	f1-score	support
	0	0.98	0.97	0.97	90
	1	0.94	0.96	0.95	53
	accuracy			0.97	143
	macro avg	0.96	0.96	0.96	143
	weighted avg	0.97	0.97	0.97	143

0.965034965034965